


Learning dynamics reveal a boundary-selective write policy

(a) Optimization discovers the high-reward regime



(b) The learned gate writes at completion

