


# Learning dynamics reveal a boundary-selective encoding policy

(a) Optimization discovers the high-reward regime



(b) The learned gate encodes at completion

